

Exp 11:

Experiment using Variational Autoencoder.

Aim:

To implement a Variational Autoencoder (VAE) in deep learning for the compression & generation of MNIST handwritten digit images.

Pseudo code:

1. Import libraries & load MNIST dataset.
2. Normalize & flatten the images
3. Define VAE model:
 - Encoder: outputs mean (μ) and log-variance
 - Reparameterization Trick: $z = \mu + \sigma * \epsilon$
 - Decoder: reconstructs image from z .
4. Define loss:
 - Reconstruction loss (MSE)
 - KL divergence loss
 - Total loss = Reconstruction + KL Divergence
5. Train model:
 - forward pass
 - compute loss
 - Backpropagate & update
6. Evaluate on test data
7. Visualize:
 - Original vs Reconstructed images.
 - Randomly generated images from latent space

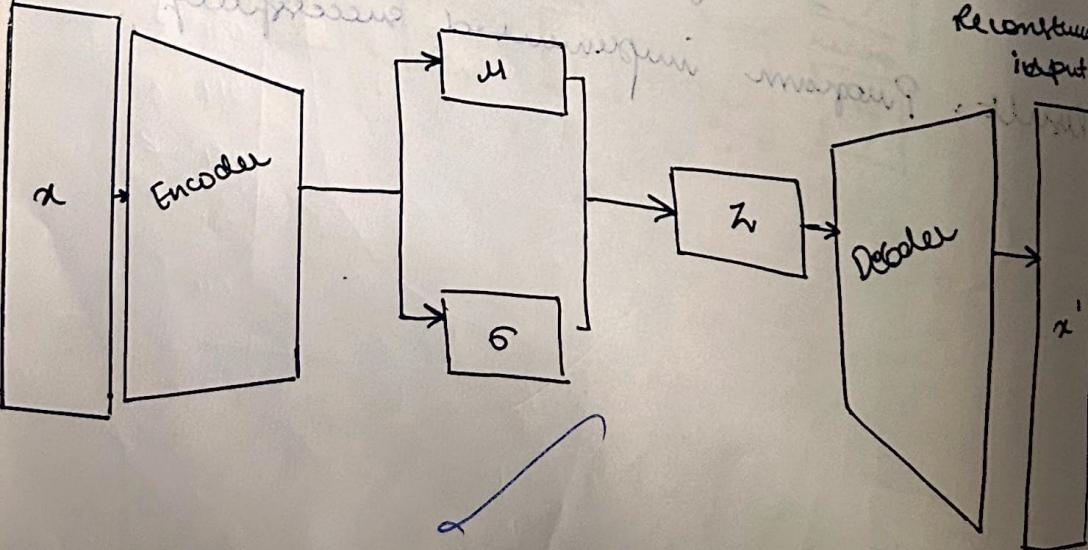
lebih besar pada saat ini dibandingkan dengan
(dalam) fungsi otak saat itu dikenal
misalkan dengan adanya teknologi
(dalam) transfer metrik (seperti teknologi
misi berfungsi setelah ada di

ini operasi

misalkan dalam
memulihkan misalkan
misalkan untuk

misalkan pula misalkan

Input.



$$z = \mu + \sigma \varepsilon$$

$$\varepsilon \sim N(0, 1)$$

Justification:

Unlike a simple Autoencoder, a VAE doesn't just learn to reconstruct input data. It learns the underlying probability distribution of the data & allows us to generate new, similar samples.

Result: Program implemented successfully.

OJ

Output:

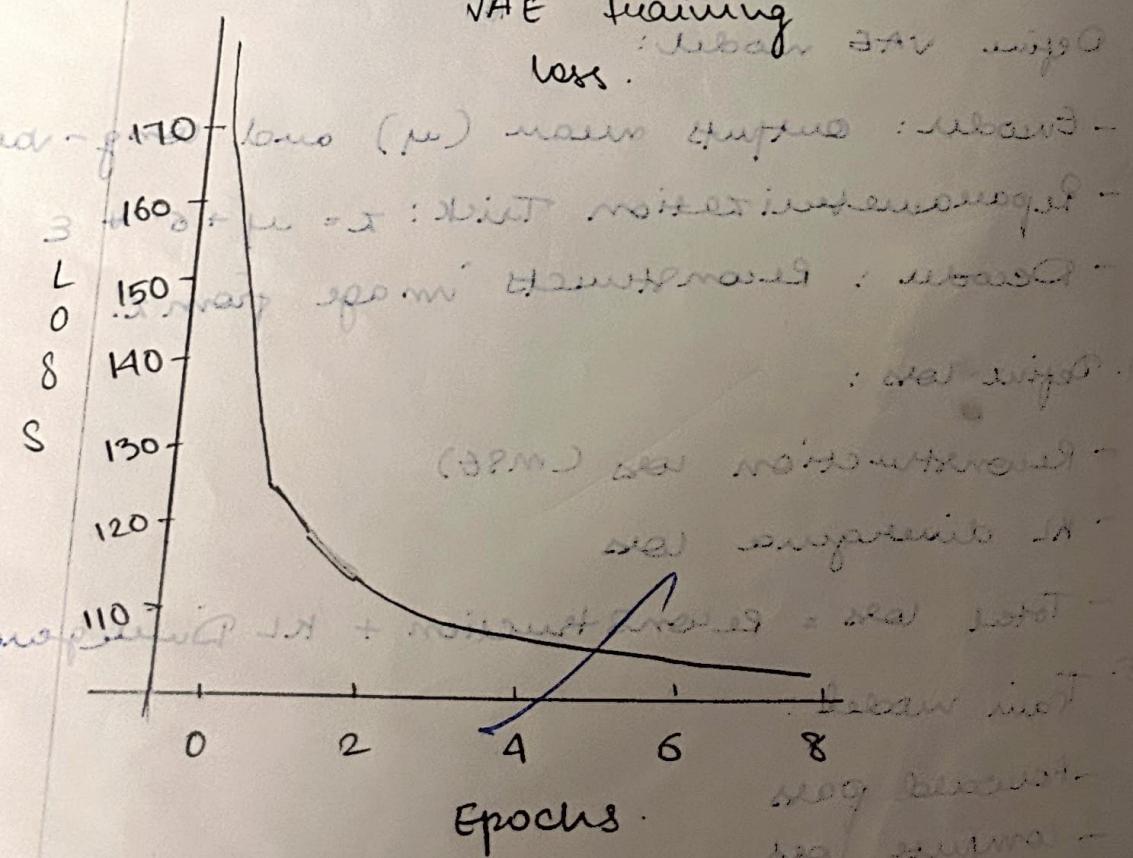
epoch [1/7], loss : 165.5449
 epoch [2/7], loss : 121.9761
 epoch [3/7], loss : 114.8512
 epoch [4/7], loss : 111.6771
 epoch [5/7], loss : 109.8832
 epoch [6/7], loss : 108.6229
 epoch [7/7], loss : 107.7510

Evaluation Metrics:

Reconstruction Accuracy : 0.9639

F1 Score : 0.8660

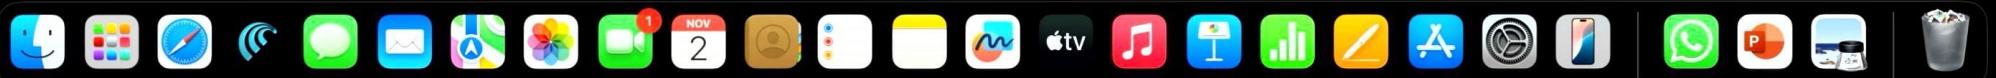
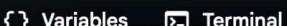
Epoch wise training & validation loss.





```
import torch
import torch.nn as nn
import torch.optim as optim
from torchvision import datasets, transforms
from torch.utils.data import DataLoader
import matplotlib.pyplot as plt
from sklearn.metrics import f1_score, accuracy_score
import numpy as np
transform = transforms.Compose([transforms.ToTensor()])
train_data = datasets.MNIST(root='./data', train=True, transform=transform, download=True)
test_data = datasets.MNIST(root='./data', train=False, transform=transform, download=True)

train_loader = DataLoader(train_data, batch_size=128, shuffle=True)
test_loader = DataLoader(test_data, batch_size=128, shuffle=False)
class VAE(nn.Module):
    def __init__(self, latent_dim=20):
        super(VAE, self).__init__()
        self.encoder = nn.Sequential(
            nn.Flatten(),
            nn.Linear(28*28, 400),
            nn.ReLU()
        )
        self.mu = nn.Linear(400, latent_dim)
        self.logvar = nn.Linear(400, latent_dim)
        self.decoder = nn.Sequential(
            nn.Linear(latent_dim, 400),
            nn.ReLU(),
            nn.Linear(400, 28*28),
            nn.Sigmoid()
        )
```



```
    nn.Linear(400, 28*28),
    nn.Sigmoid()
)

def reparameterize(self, mu, logvar):
    std = torch.exp(0.5 * logvar)
    eps = torch.randn_like(std)
    return mu + eps * std

def forward(self, x):
    encoded = self.encoder(x)
    mu = self.mu(encoded)
    logvar = self.logvar(encoded)
    z = self.reparameterize(mu, logvar)
    decoded = self.decoder(z)
    return decoded, mu, logvar

def vae_loss(recon_x, x, mu, logvar):
    BCE = nn.functional.binary_cross_entropy(recon_x, x.view(-1, 28*28), reduction='sum')
    KLD = -0.5 * torch.sum(1 + logvar - mu.pow(2) - logvar.exp())
    return BCE + KLD

device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
model = VAE().to(device)
optimizer = optim.Adam(model.parameters(), lr=0.001)
train_losses = []
for epoch in range(10):
    model.train()
    total_loss = 0

    for batch_idx, (data, _) in enumerate(train_loader):
        data = data.to(device)
        optimizer.zero_grad()
        recon, mu, logvar = model(data)
```

Untitled8.ipynb ☆ Saving failed since 2:39 PM

File Edit View Insert Runtime Tools Help

Commands + Code + Text ▶ Run all ⌘

Connecting A

```
total_loss = 0

for batch_idx, (data, _) in enumerate(train_loader):
    data = data.to(device)
    optimizer.zero_grad()
    recon, mu, logvar = model(data)
    loss = vae_loss(recon, data, mu, logvar)
    loss.backward()
    total_loss += loss.item()
    optimizer.step()

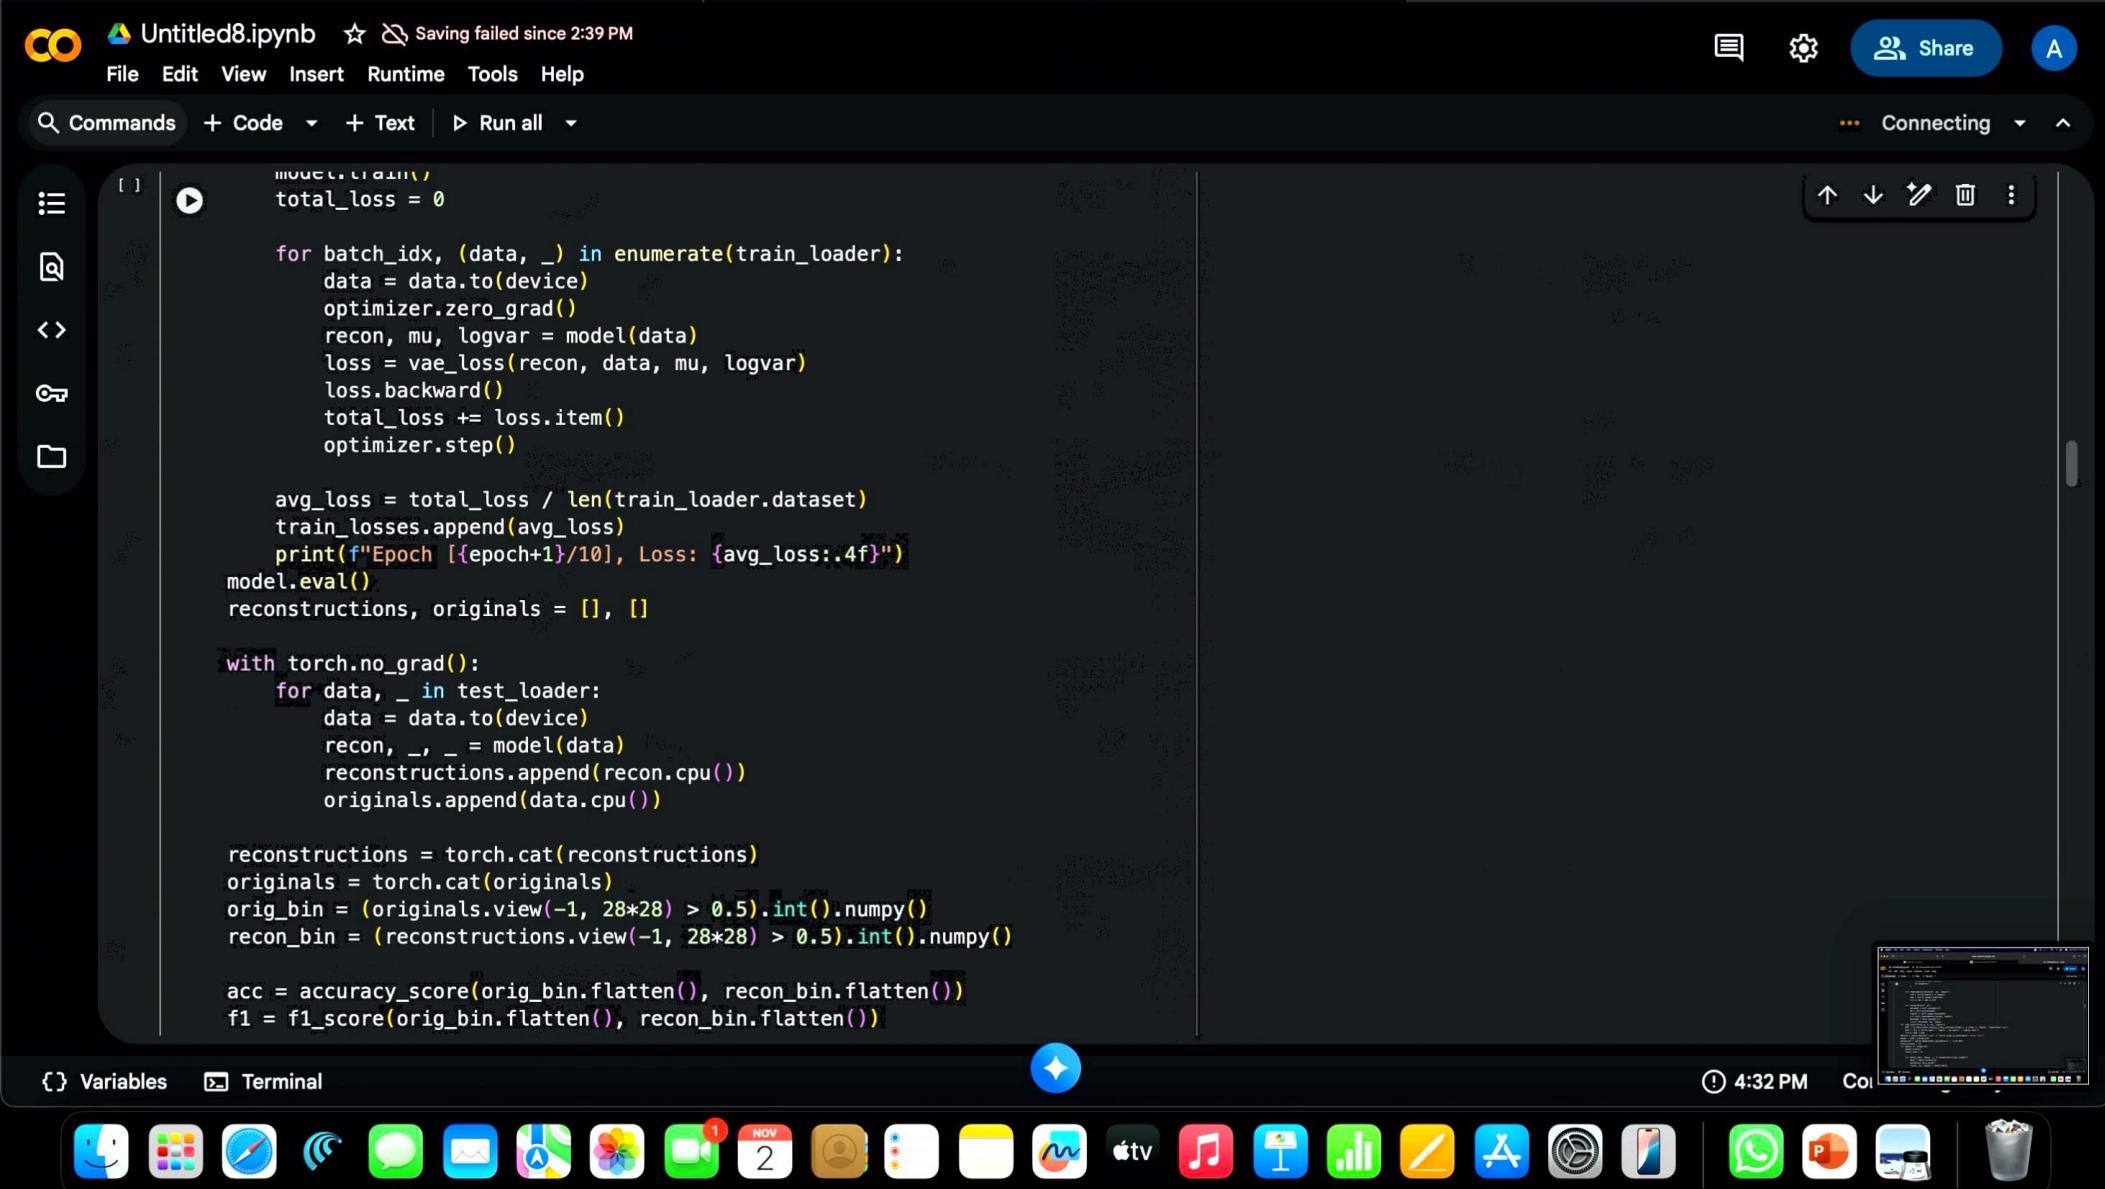
    avg_loss = total_loss / len(train_loader.dataset)
    train_losses.append(avg_loss)
    print(f"Epoch [{epoch+1}/10], Loss: {avg_loss:.4f}")
model.eval()
reconstructions, originals = [], []

with torch.no_grad():
    for data, _ in test_loader:
        data = data.to(device)
        recon, _, _ = model(data)
        reconstructions.append(recon.cpu())
        originals.append(data.cpu())

reconstructions = torch.cat(reconstructions)
originals = torch.cat(originals)
orig_bin = (originals.view(-1, 28*28) > 0.5).int().numpy()
recon_bin = (reconstructions.view(-1, 28*28) > 0.5).int().numpy()

acc = accuracy_score(orig_bin.flatten(), recon_bin.flatten())
f1 = f1_score(orig_bin.flatten(), recon_bin.flatten())
```

Variables Terminal ⌘ 4:32 PM Cor





```
[ ] acc = accuracy_score(orig_bin.flatten(), recon_bin.flatten())
f1 = f1_score(orig_bin.flatten(), recon_bin.flatten())

print("\n\ufe0f Evaluation Metrics:")
print(f"Reconstruction Accuracy: {acc:.4f}")
print(f"F1 Score: {f1:.4f}")
plt.figure(figsize=(6,4))
plt.plot(train_losses, label='Train Loss')
plt.title("VAE Training Loss")
plt.xlabel("Epochs")
plt.ylabel("Loss")
plt.legend()
plt.show()
n = 8
plt.figure(figsize=(12,4))
for i in range(n):
    plt.subplot(2, n, i+1)
    plt.imshow(originals[i][0], cmap="gray")
    plt.axis("off")
    plt.subplot(2, n, n+i+1)
    plt.imshow(reconstructions[i].view(28,28), cmap="gray")
    plt.axis("off")

plt.suptitle("Top: Original | Bottom: Reconstructed", fontsize=14)
plt.show()
```

```
100% [██████] 9.91M/9.91M [00:00<00:00, 60.3MB/s]
100% [██████] 28.9k/28.9k [00:00<00:00, 1.88MB/s]
100% [██████] 1.65M/1.65M [00:00<00:00, 15.1MB/s]
100% [██████] 4.54k/4.54k [00:00<00:00, 7.15MB/s]
Epoch 1/101. Loss: 165.5449
```



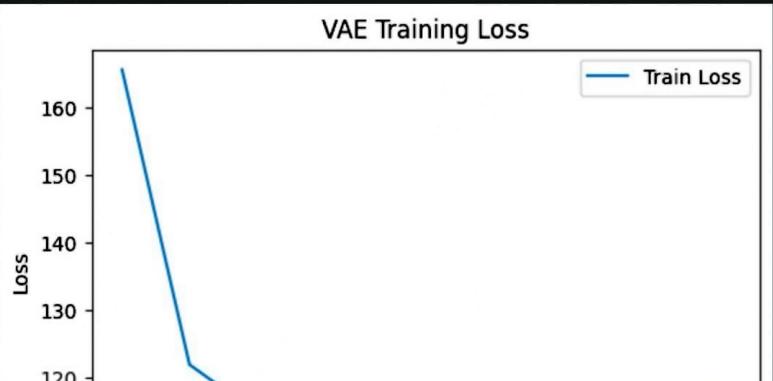
Commands + Code + Text | ▶ Run all

... Connecting ▾



```
100%|██████████| 9.91M/9.91M [00:00<00:00, 60.3MB/s]
100%|██████████| 28.9k/28.9k [00:00<00:00, 1.88MB/s]
100%|██████████| 1.65M/1.65M [00:00<00:00, 15.1MB/s]
100%|██████████| 4.54k/4.54k [00:00<00:00, 7.15MB/s]
Epoch [1/10], Loss: 165.5449
Epoch [2/10], Loss: 121.9764
Epoch [3/10], Loss: 114.8512
Epoch [4/10], Loss: 111.6771
Epoch [5/10], Loss: 109.8832
Epoch [6/10], Loss: 108.6224
Epoch [7/10], Loss: 107.7510
Epoch [8/10], Loss: 107.0746
Epoch [9/10], Loss: 106.5901
Epoch [10/10], Loss: 106.1244
```

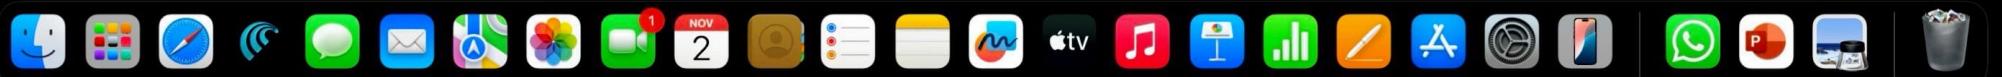
Evaluation Metrics:
Reconstruction Accuracy: 0.9639
F1 Score: 0.8660

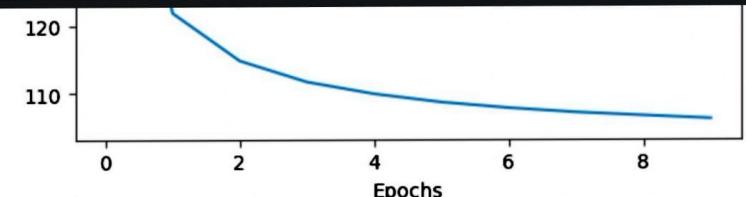


Variables Terminal



① 4:32 PM Col





Top: Original | Bottom: Reconstructed



```
▶ import torch
    import torch.nn as nn
    import torch.optim as optim
    from torchvision import datasets, transforms
```

