

tp114. Implement a pre-trained CNN model as a feature extractor using transfer learning.

Aim:

Implementing a pre-trained CNN model as a feature extractor using transfer learning.

Pseudo code:

1. Import necessary libraries (torch, torchvision, matplotlib, etc.)
2. Load the pre-trained CNN model (eg: ResNet50 or VGG16) from torchvision.models
3. Freeze the convolutional base to retain pre-trained weights.
4. Replace the final classification layer w/ a new one suitable for the new dataset's no. of classes.
5. Load the dataset and apply image transformations (resize, normalize)
6. Extract features using the frozen model.
7. Train only the newly added layers.
8. Evaluate model performance using accuracy & f1-score.
9. Visualize training / validation accuracy & loss.

Justification:-

CNNs are deep learning architectures widely used for img. classification & recognition:


They consist of:

- conventional layers for feature extraction
- Pooling layers for downsampling

- fully connected layer for classification

~~Result~~: Program implemented successfully,

BT



TP colab.research.google.com

clean coal technology diagrams - Google Search

SRM Student Portal

Untitled8.ipynb - Colab

21AIC301J- Lab AI-A

Untitled8.ipynb Saving failed since 7:38 PM

File Edit View Insert Runtime Tools Help

Commands + Code + Text Run all

Share RAM Disk

Epoch

```
[2] import torch
import torch.nn as nn
import torchvision.models as models
from torch.utils.data import DataLoader, TensorDataset
from sklearn.metrics import accuracy_score, f1_score, recall_score
import matplotlib.pyplot as plt
import numpy as np
class PretrainedCNNFeatureExtractor(nn.Module):
    def __init__(self, output_dim=2):
        super(PretrainedCNNFeatureExtractor, self).__init__()
        self.resnet = models.resnet18(pretrained=True)
        for param in self.resnet.parameters():
            param.requires_grad = False
        num_features = self.resnet.fc.in_features
        self.resnet.fc = nn.Linear(num_features, output_dim)
    def forward(self, x):
        return self.resnet(x)
np.random.seed(0)
X_train = np.random.rand(100, 3, 64, 64).astype(np.float32)
y_train = np.random.randint(0, 2, 100)
X_test = np.random.rand(30, 3, 64, 64).astype(np.float32)
y_test = np.random.randint(0, 2, 30)
train_data = TensorDataset(torch.from_numpy(X_train), torch.from_numpy(y_train))
test_data = TensorDataset(torch.from_numpy(X_test), torch.from_numpy(y_test))

train_loader = DataLoader(train_data, batch_size=16, shuffle=True)
test_loader = DataLoader(test_data, batch_size=10, shuffle=False)
model = PretrainedCNNFeatureExtractor(output_dim=2)
criterion = nn.CrossEntropyLoss()
```

Variables Terminal ✓ 7:43 PM Python 3

4 NOV

clean coal technology diagrams - Google Search

SRM Student Portal

Untitled8.ipynb - Colab

21AIC301J- Lab Ai-A

Untitled8.ipynb Saving failed since 7:38 PM

File Edit View Insert Runtime Tools Help

Commands + Code + Text Run all

RAM Disk

2 train\_loader = DataLoader(train\_data, batch\_size=16, shuffle=True)  
test\_loader = DataLoader(test\_data, batch\_size=10, shuffle=False)  
model = PretrainedCNNFeatureExtractor(output\_dim=2)  
criterion = nn.CrossEntropyLoss()  
optimizer = torch.optim.Adam(model.resnet.fc.parameters(), lr=0.001)  
num\_epochs = 10  
train\_losses = []  
for epoch in range(num\_epochs):  
 model.train()  
 running\_loss = 0.0  
 for inputs, labels in train\_loader:  
 optimizer.zero\_grad()  
 outputs = model(inputs)  
 loss = criterion(outputs, labels)  
 loss.backward()  
 optimizer.step()  
 running\_loss += loss.item() \* inputs.size(0)  
 epoch\_loss = running\_loss / len(train\_loader.dataset)  
 train\_losses.append(epoch\_loss)  
 print(f'Epoch {epoch+1}/{num\_epochs}, Loss: {epoch\_loss:.4f}')  
model.eval()  
all\_preds, all\_labels = [], []  
with torch.no\_grad():  
 for inputs, labels in test\_loader:  
 outputs = model(inputs)  
 \_, preds = torch.max(outputs, 1)  
 all\_preds.extend(preds.numpy())  
 all\_labels.extend(labels.numpy())  
accuracy = accuracy\_score(all\_labels, all\_preds)  
f1 = f1\_score(all\_labels, all\_preds)  
recall = recall\_score(all\_labels, all\_preds)  
print(f'Accuracy: {accuracy:.4f}')

Variables Terminal

7:43

clean coal technology diagrams - Google Search

SRM Student Portal

Untitled8.ipynb - Colab

21AIC301J- Lab AI-A

Untitled8.ipynb Saving failed since 7:38 PM

File Edit View Insert Runtime Tools Help

Commands + Code + Text ▶ Run all

RAM Disk

[2]

```
_preds = torch.max(outputs, 1)
all_preds.extend(preds.numpy())
all_labels.extend(labels.numpy())
accuracy = accuracy_score(all_labels, all_preds)
f1 = f1_score(all_labels, all_preds)
recall = recall_score(all_labels, all_preds)
print(f'Accuracy: {accuracy:.4f}')
print(f'F1 Score: {f1:.4f}')
print(f'Recall: {recall:.4f}')
plt.plot(train_losses, label='Train Loss')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.title('Training Loss Curve')
plt.legend()
plt.show()
```

/usr/local/lib/python3.12/dist-packages/torchvision/models/\_utils.py:208: UserWarning: The parameter 'pretrained' is deprecated since 0.13 and may be removed in a future release.
 warnings.warn(
/usr/local/lib/python3.12/dist-packages/torchvision/models/\_utils.py:223: UserWarning: Arguments other than a weight enum or `None` for 'weights' are deprecated.
 warnings.warn(msg)
Downloading: "<https://download.pytorch.org/models/resnet18-f37072fd.pth>" to /root/.cache/torch/hub/checkpoints/resnet18-f37072fd.pth
100% [██████████] 44.7M/44.7M [00:00<00:00, 127MB/s]
Epoch 1/10, Loss: 0.7812
Epoch 2/10, Loss: 0.6875
Epoch 3/10, Loss: 0.6530
Epoch 4/10, Loss: 0.5271
Epoch 5/10, Loss: 0.5459
Epoch 6/10, Loss: 0.5682
Epoch 7/10, Loss: 0.5525
Epoch 8/10, Loss: 0.5192
Epoch 9/10, Loss: 0.4761

Variables Terminal ✓ 7:43

clean coal technology diagrams - Google Search

SRM Student Portal

Untitled8.ipynb - Colab

21AIC301J- Lab Ai-A

Untitled8.ipynb Saving...

File Edit View Insert Runtime Tools Help

Commands + Code + Text Run all

RAM Disk

Epoch 18/20, Loss: 0.6059  
Epoch 19/20, Loss: 0.5769  
Epoch 20/20, Loss: 0.5681  
Accuracy: 0.5333  
F1 Score: 0.5333  
Recall: 0.6154

Training Loss Curve

Epoch	Loss
0.0	0.69
2.5	0.685
5.0	0.68
7.5	0.67
10.0	0.64
12.5	0.61
15.0	0.59
17.5	0.575

[2] import torch

Variables Terminal ✓ 7:43

4

W