

Exp:4

BUILD A SIMPLE FEED FORWARD NEURAL NETWORK
TO RECOGNIZE HANDWRITTEN CHARACTER.

14.08.25

Aim:

To build a simple feed forward neural network
to recognize hand written character.


Pseudo
code:

1. Import necessary libraries.
→ (numpy, tensorflow / keras, matplotlib)
2. Load dataset
→ (e.g.: MNIST / EMNIST of handwritten character)
3. Preprocess data
 - a. Normalize pixel values to range [0, 1]
 - b. Flatten each image into a 1D vector.
 - c. Convert labels into one-hot encoding.
4. Define FFNN architecture.
 - a. Input layer : size = no. of pixels per image
 - b. Hidden layers : fully connected neurons with activation (ReLU)
 - c. Output layer : size = no. of character classes, activation (Softmax).
5. Compile the model
 - a. loss function : Categorical Cross - Entropy.
 - b. Optimizers : SGD / Adam.
 - c. Metrics : Accuracy, F1 Score, recall.
6. Train the model.
 - a. fit training data for N epochs.
 - b. Validate on test data.

7. Evaluate model performance.
- a. Test accuracy
 - b. Confusion matrix and classification report
8. Display sample predictions w/ true labels.

Justification:

1. Choice of model. (FFNN):
→ Represents the foundational arch. of Deep learning
2. Learning outcome:
 - Understanding of NN layers & activation func.
 - Practical exposure to training, validation & testing in supervised learning



Observation

- i. Training & Validation Accuracy:
 - Training accuracy $\approx 98.85\%$ after 20 epochs.
 - Validation accuracy $\approx 98.47\%$.
- fitted well to unseen data, indicating a good model generalization.

Accuracy vs. number of training epochs
(dotted line: validation accuracy, solid line: training accuracy)

Labels with number of training (train : val) -

(1, 0) signs of overfitting although no token at 0 has been seen which probably has no other label than 1

unseen words with signs of being far from the type token with similar meaning.

```
1 import numpy as np
2 import tensorflow as tf
3 from tensorflow.keras.models import Sequential
4 from tensorflow.keras.layers import Dense, Dropout
5 from tensorflow.keras.callbacks import EarlyStopping
6 from tensorflow.keras.datasets import mnist
7 (X_train, y_train), (X_test, y_test) = mnist.load_data()
8 X_train = X_train.reshape(-1, 28*28).astype("float32") / 255.0
9 X_test = X_test.reshape(-1, 28*28).astype("float32") / 255.0
10 print("Training set shape:", X_train.shape, y_train.shape)
11 print("Test set shape:", X_test.shape, y_test.shape)
12 model = Sequential([
13     Dense(256, activation='relu', input_shape=(784,)),
14     Dropout(0.3),
15     Dense(128, activation='relu'),
16     Dropout(0.2),
17     Dense(64, activation='relu'),
18     Dense(10, activation='softmax')
19 ])
20 model.compile(
21     optimizer='adam',
22     loss='sparse_categorical_crossentropy',
23     metrics=['accuracy']
24 )
25 early_stop = EarlyStopping(monitor='val_loss', patience=3, restore_best_weights=True)
26 history = model.fit(
27     X_train, y_train,
28     epochs=20,
29     batch_size=64,
30     validation_split=0.1,
```

```
Week2.py Week3.py Untitled.ipynb untitled.txt week4.py jupyter lab231104.ipynb
9 X_test = X_test.reshape(-1, 28*28).astype("float32") / 255.0
10 print("Training set shape:", X_train.shape, y_train.shape)
11 print("Test set shape:", X_test.shape, y_test.shape)
12 model = Sequential([
13     Dense(256, activation='relu', input_shape=(784,)),
14     Dropout(0.3),
15     Dense(128, activation='relu'),
16     Dropout(0.2),
17     Dense(64, activation='relu'),
18     Dense(10, activation='softmax')
19 ])
20 model.compile(
21     optimizer='adam',
22     loss='sparse_categorical_crossentropy',
23     metrics=['accuracy']
24 )
25 early_stop = EarlyStopping(monitor='val_loss', patience=3, restore_best_weights=True)
26 history = model.fit(
27     X_train, y_train,
28     epochs=20,
29     batch_size=64,
30     validation_split=0.1,
31     callbacks=[early_stop],
32     verbose=2
33 )
34 loss, accuracy = model.evaluate(X_test, y_test)
35 print(f"\nTest Accuracy: {accuracy * 100:.2f}%")
36 model.save("mnist_ffnn_model.h5")
37 print("Model saved as mnist_ffnn_model.h5")
38
```

To enable the following instructions: AVX2 FMA, in other operations, rebuild TensorFlow with the appropriate compiler flags.

Training set shape: (60000, 784) (60000,)

Test set shape: (10000, 784) (10000,)

/home/jupyter-ra2311047010012/.local/lib/python3.10/site-packages/keras/src/layers/core/dense.py:92: UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.

super().__init__(activity_regularizer=activity_regularizer, **kwargs)

Epoch 1/20

844/844 - 7s - 9ms/step - accuracy: 0.8932 - loss: 0.3489 - val_accuracy: 0.9667 - val_loss: 0.1079

Epoch 2/20

844/844 - 5s - 5ms/step - accuracy: 0.9527 - loss: 0.1545 - val_accuracy: 0.9763 - val_loss: 0.0787

Epoch 3/20

844/844 - 5s - 6ms/step - accuracy: 0.9623 - loss: 0.1227 - val_accuracy: 0.9785 - val_loss: 0.0694

Epoch 4/20

844/844 - 5s - 6ms/step - accuracy: 0.9701 - loss: 0.0983 - val_accuracy: 0.9802 - val_loss: 0.0684

Epoch 5/20

844/844 - 5s - 6ms/step - accuracy: 0.9723 - loss: 0.0895 - val_accuracy: 0.9805 - val_loss: 0.0680

Epoch 6/20

844/844 - 5s - 6ms/step - accuracy: 0.9756 - loss: 0.0773 - val_accuracy: 0.9813 - val_loss: 0.0650

Epoch 7/20

844/844 - 5s - 6ms/step - accuracy: 0.9768 - loss: 0.0735 - val_accuracy: 0.9812 - val_loss: 0.0616

Epoch 8/20

844/844 - 5s - 6ms/step - accuracy: 0.9792 - loss: 0.0650 - val_accuracy: 0.9797 - val_loss: 0.0754

Epoch 9/20

844/844 - 5s - 6ms/step - accuracy: 0.9813 - loss: 0.0591 - val_accuracy: 0.9830 - val_loss: 0.0598

Epoch 10/20

844/844 - 5s - 6ms/step - accuracy: 0.9816 - loss: 0.0585 - val_accuracy: 0.9833 - val_loss: 0.0569

Epoch 11/20

844/844 - 5s - 6ms/step - accuracy: 0.9837 - loss: 0.0518 - val_accuracy: 0.9810 - val_loss: 0.0730

Epoch 12/20

844/844 - 5s - 5ms/step - accuracy: 0.9832 - loss: 0.0520 - val_accuracy: 0.9807 - val_loss: 0.0668

Epoch 13/20

844/844 - 5s - 6ms/step - accuracy: 0.9849 - loss: 0.0475 - val_accuracy: 0.9830 - val_loss: 0.0649

313/313 1s 3ms/step - accuracy: 0.9814 - loss: 0.0648