

16.09.25

BUILD A CNN MODEL TO CLASSIFY CAT & DOG IMAGE.

Aim:

To implement CNN to build a CNN model to classify a cat & dog image.

PSEUDO
CODE:

- Initialize parameters
 - Define input of shape = $(64, 64, 3)$ → image resolution chosen for uniformity.
 - Define no. of classes = 2 → binary classification (Cat, Dog)
 - Set batch size, learning rate & epochs.
(These hyperparameters control how the model learns).
- load dataset
 - Import cat & dog image from dataset
 - Preprocess: resize all images to same dimension.
 - Normalize pixels values [0, 1] to stabilize training.
 - Split into training set (for learning) & validation set (for performance check)
- Data Augmentation
 - Apply random flips, rotation, zooms & shifts
(increases diversity of training images & prevents overfitting)
- Construct CNN architecture
 - Convolution layer: Applies filters to capture local spatial features (e.g. edges, corners, textures of cats/dogs)
 - Activation function (ReLU): Introduces non-linearity, allowing complex feature learning.
 - Pooling layer: Reduces spatial dimension.
 - Repeat convolution + pooling block to learn higher-level features.

- compile model.

- optimizer = adam (adaptive learning, efficient convergence)

- loss function = Binary cross entropy.

- Evaluation metric = Accuracy.

- Train model.

- feed training set in batches.

- validate on validation set after each epoch.

- repeat for defined no. of epochs.

- Evaluate model.

- calculate accuracy & loss on validation set.

- Prediction.

- for a new image:

- Resize to (64, 64)

- Normalize pixel values.

- Pass through trained CNN.

- If output $< 0.5 \rightarrow$ Cat

- \rightarrow Dog.

Justification.


- Why CNN?

- Images have spatial patterns (edges, colors, textures, shapes).

- Specifically designed to capture these local patterns by applying filters (kernels).


- Why multiple convolution + pooling layers?

CNN ARCHITECTURE DIAGRAM.



OBSERVATION

loss of loss function at training pass & loss is given by CNN training loss curve (cats vs dogs).



Epochs along x-axis -
initials of (1, 0)

Table Epochs training validation loss notes.
Accuracy improving

1	65%	60%	0.82	Model still learning
5	88%	82%	0.4	Overfitting not yet seen
10	92%	85%	0.3	Good gradation.

Untitled0.ipynb

File Edit View Insert Runtime Tools Help

Connect 14

Commands + Code + Text ▶ Run all

```
model = SimpleCNN(img_size=IMG_SIZE).to(device)
print(model)

SimpleCNN(
    (features): Sequential(
        (0): Conv2d(3, 32, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (1): ReLU()
        (2): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
        (3): Conv2d(32, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (4): ReLU()
        (5): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
        (6): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (7): ReLU()
        (8): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
    )
    (classifier): Sequential(
        (0): Flatten(start_dim=1, end_dim=-1)
        (1): Linear(in_features=41472, out_features=512, bias=True)
        (2): ReLU()
        (3): Dropout(p=0.5, inplace=False)
        (4): Linear(in_features=512, out_features=1, bias=True)
        (5): Sigmoid()
    )
)

# Loss and optimizer
criterion = nn.BCELoss()
optimizer = optim.Adam(model.parameters(), lr=0.001)

#Training loop
for epoch in range(NUM_EPOCHS):
    model.train()
    running_loss = 0.0
```

() Variables () Terminal

Untitled0.ipynb

File Edit View Insert Runtime Tools Help

Connect 14

Commands + Code + Text ▶ Run all

```
optimizer.zero_grad()
outputs = model(inputs)
loss = criterion(outputs, labels)
loss.backward()
optimizer.step()

running_loss += loss.item() * inputs.size(0)
preds = (outputs > 0.5).float()
running_corrects += torch.sum(preds == labels)

epoch_loss = running_loss / len(train_dataset)
epoch_acc = running_corrects.double() / len(train_dataset)

print(f'Epoch {epoch+1}/{NUM_EPOCHS} - Loss: {epoch_loss:.4f} Acc: {epoch_acc:.4f}')

Epoch 1/10 - Loss: 0.7142 Acc: 0.5440
Epoch 2/10 - Loss: 0.6549 Acc: 0.6070
Epoch 3/10 - Loss: 0.6254 Acc: 0.6640
Epoch 4/10 - Loss: 0.5928 Acc: 0.6765
Epoch 5/10 - Loss: 0.5811 Acc: 0.6900
Epoch 6/10 - Loss: 0.5511 Acc: 0.7200
Epoch 7/10 - Loss: 0.5089 Acc: 0.7570
Epoch 8/10 - Loss: 0.4831 Acc: 0.7670
Epoch 9/10 - Loss: 0.4617 Acc: 0.7870
Epoch 10/10 - Loss: 0.4135 Acc: 0.8075

validation
model.eval()
val_corrects = 0

with torch.no_grad():
    for inputs, labels in val_loader:
        inputs = inputs.to(device)
```

() Variables () Terminal